

Feasibility Analysis

MIMIC-IV QA Benchmark for Coverage vs. Dispersion

Contents

1	Introduction	2
I	MIMIC-IV Dataset	2
2	MIMIC-IV-Note	3
2.1	Discharge Summaries	3
2.2	Radiology Reports	4
2.3	Considerations	5
2.3.1	Patient-specific vs. Generalizable data	5
3	MIMIC-IV Structured Data	5
3.1	Hosp Module	5
3.2	ICU Module	6
II	QA Dataset Generation Pipeline	7
4	Pipeline draft	7
4.1	Phases	7
	Phase 1: Corpus Construction	7
	Phase 2: Embeddings Generation	7
	Phase 3: Query Generation	8
	Phase 4: Evaluation	9
4.2	Does Coverage Pull in Noise?	10
4.3	Computational Feasibility	10
5	LLM Usage and Data Use Agreement	10
6	Embeddings: considerations	11
6.1	Contextual Embeddings	11
6.2	Chunking vs. Whole-Document Embedding	11
6.3	Embedding Models for medical domain	11
7	Cluster Validation	11
7.1	Pipeline draft	11
7.2	Is Full Clustering Necessary?	12
7.3	High-Dimensional Clustering: Caveats	12
	Appendix	13
A	Available Analytics from mimic-code	13

1 Introduction

The experiment’s goal is that *coverage* (facility-location) outperforms *dispersion* (MMR, gMMR, FPS) when the candidate pool of retrieved chunks has a multi-cluster structure. Other QA datasets exist (HotpotQA, 2WikiMultiHopQA, MuSiQue), but these benchmarks have fixed candidate pools that we cannot control.

Ideal Cluster Structure for coverage

1. **Redundant central cluster** - dominant topic, many similar chunks. Top- k over-samples from here. Coverage naturally avoids redundancy.
2. **Complementary cluster(s)** - other relevant aspects of the query. Coverage ensures at least one representative from each.
3. **Niche cluster** - small but informative, not an outlier. Coverage includes it because facility-location rewards covering every region of D .
4. **Outliers** - informative or noisy. Dispersion favors them (maximizes spread); coverage ignores them unless they represent a dense region.

Proposal: build a **custom QA benchmark from MIMIC-IV** where we can (a) verify that candidate pools have the desired cluster structure, (b) control pool composition, and (c) design multi-aspect queries that require many sources.

Part I

MIMIC-IV Dataset

The database is split across two datasets (important CSVs highlighted in blue):

- **MIMIC-IV** $\Rightarrow$ tabular clinical data recording every diagnosis, procedure, and medication for each hospitalization. This can give us **ground truth** for evaluation. There are two submodules:
 - **hosp**: patient tracking ([patients](#), [admissions](#), [transfers](#)), diagnoses ([diagnoses_icd](#), [d_icd_diagnoses](#), [procedures_icd](#), [drgcodes](#)), labs ([labevents](#)), medications ([prescriptions](#), [emar](#)).
 - **icu**: bedside data ([chartevents](#), [inputevents](#), [outputevents](#), [procedureevents](#)).
- **MIMIC-IV-Note** $\Rightarrow$ free-text clinical notes
 - [discharge](#): discharge summaries. Long, narrative clinical documents written by doctors when a patient leaves the hospital.
 - [radiology](#): radiology reports.

For our purposes, the critical tables are:

- [discharge](#) (in MIMIC-IV-Note): $\sim$ 331K discharge summaries.
- [diagnoses_icd](#): join table linking ICD-9/ICD-10 (International Classification of Diseases) codes to subject hospitalizations, with approximate importance ordering ([seq_num](#)). A typical admission has 5-20 ICD codes, often spanning multiple body systems. Definitions for the codes themselves are in the [d_icd_diagnoses](#) table
- [patients](#): gender, [anchor_age](#), date of death.
- [admissions](#): admission type, insurance, language, marital status, race, discharge location.

All tables link via `subject_id` (patient) and `hadm_id` (hospitalization). The dataset has 331,793 discharge summaries from 145,914 patients. All protected health information (names, dates, locations) has been removed or replaced with placeholders (---).

Some analytics about demographics:

	Hospital admissions	ICU admissions
Number of stays	546,028	94,458
Unique patients	223,452	65,366
Age, mean (SD)	59.2 (19.1)	64.7 (16.9)
Female Administrative Gender, n (%)	284,097 (52.0)	37,714 (44.2)
Hospital length of stay, mean (SD)	4.8 (7.2)	3.6 (5.4)
In-hospital mortality, n (%)	11,801 (2.2)	9,475 (11.1)
One year mortality, n (%)	64,934 (11.9)	14,951 (17.5)
Data period	2008–2022	

2 MIMIC-IV-Note

The notes are of two kinds: **Discharge Summaries** and **Radiology Reports**.

2.1 Discharge Summaries

The structure of the notes seems consistent, with standardized sections:

Section	Content	QA Value
Brief Hospital Course	Clinical reasoning per problem	Highest
History of Present Illness	Symptom narrative, temporal progression	High
Pertinent Results	Labs, imaging, micro interpretations	High
Discharge Diagnosis	Diagnostic labels	High (linkable to ICD)
Chief Complaint	1–2 lines, reason for admission	High only as contextual embedding metadata
Past Medical History	Comorbidities list	Medium
Physical Exam	Admission + discharge findings	Medium
Discharge Medications	Treatment protocols with dosing	Medium
Major Surgical or Invasive Procedure	Procedures performed during stay	Low
Allergies	-	Low
Family History	-	Low
Medications on Admission	-	Low
Discharge Disposition	Destination (home, SNF, rehab, ...)	Low
Discharge Condition	Mental status, consciousness, activity	Low
Discharge Instructions	Patient-facing plain-language summary	Low
Service	Medicine, Surgery, ...	None
Attending	Attending physician	None (deidentified)
Social History	Lifestyle, substance use	None (deidentified)
Followup Instructions	Outpatient follow-up appointments	None (deidentified)

Brief Hospital Course. This section uses `#` markers to delimit individual clinical problems (e.g., `# Stroke`, `# Hypertension`, `# DVT`). Each sub-section contains *clinical reasoning*: differential diagnoses, treatment rationale, management decisions with justifications. This is where generalizable medical knowledge is embedded within patient-specific narratives.

However, the `#`-header format is not universal. There are two patterns:

1. **Complex medicine admissions** (the interesting cases): the Hospital Course is organized as a problem list with # headers, one block per active clinical issue. Example: a cirrhosis patient’s Hospital Course has # ASCITES, # HEPATIC ENCEPHALOPATHY, # HYPONATREMIA, # HIV, # COPD (5–8 blocks, each 50–200 tokens). Sometimes there are additional sub-headers like CHRONIC ISSUES or ACUTE/ACTIVE ISSUES grouping the problems. The formatting is inconsistent: # Ascites, # ASCITES, #NSTEMI (with or without space after #, mixed casing).
2. **Simple or surgical admissions** (single-problem cases): the Hospital Course is a single narrative paragraph with *no* # headers. This includes orthopedic (hip fracture), urological (nephrectomy), and straightforward medicine cases (dysphagia workup). These are analogous to single-hop queries. There is only one information peak, so coverage should add nothing.

The TRANSITIONAL ISSUES block. This (sometimes misspelled as TRANISTIONAL ISSUES) appears at the end of many Hospital Courses and contains follow-up action items (“recheck electrolytes”, “follow up with hepatology”), not clinical narratives. These are task-oriented and do not describe clinical problems.

Deidentification artifacts. The ___ placeholder is everywhere in dates, names, ages, and locations. Brief Hospital Course and Physical Exam can contain few of them but largely are usable. Ages are recoverable from `anchor_age` in the `patients` table.

2.2 Radiology Reports

The `radiology` table in MIMIC-IV-Note contains ~2.3M radiologist reports spanning x-ray, CT, MRI, and ultrasound. Reports follow a structured format with sections: **Examination**, **Indication**, **Technique**, **Comparison**, **Findings**, and **Impression**. The associated `radiology_detail` table provides exam codes and exam names (e.g., CT HEAD W/O CONTRAST, LIVER OR GALLBLADDER US) that can be used for filtering.

Content quality is dual:

- **Formulaic reports** (routine chest x-rays, follow-up films): short, low information content. Typical impression: *“No acute cardiopulmonary abnormality.”*
- **Detailed diagnostic reports** (CT head, CT abdomen with contrast, MRI brain, CT angiography): substantial findings sections with organ-by-organ descriptions, measurements, and clinical correlation. Example from a CT head: *“Subcortical and periventricular white matter hypodensities consistent with chronic small vessel ischemic disease. No evidence of acute intracranial hemorrhage, edema, mass effect, or large territorial infarction.”*

Potential value for coverage vs. dispersion. For imaging-dependent conditions (stroke, PE, pneumonia), radiology reports provide a **different semantic perspective** from discharge notes: the discharge note describes clinical management (“tPA administered within 3h window”), while the radiology report describes anatomical findings (“acute hypodensity in L MCA territory, no midline shift”). This seems good for coverage.

Redundancy with discharge notes. The “Pertinent Results” section of discharge summaries often summarizes key imaging findings. Radiology reports may not add unique facets beyond what is already captured there.

I think Discharge Notes are better as a starting point and this can be used later if needed.

2.3 Considerations

2.3.1 Patient-specific vs. Generalizable data

Discharge notes are patient-specific narratives, not medical textbooks. A query like “*what are the general symptoms of a stroke?*” has no ground truth in this corpus: no individual chunk states textbook symptomatology. Embedding similarity will match chunks that *mention* stroke, but the question is unanswerable from patient notes alone.

The queries should ask about **clinical management patterns**: things that are described in discharge notes (the Brief Hospital Course says what was done and why). A more suitable query:

“How is pneumonia managed in an elderly patient with chronic kidney failure?”

This works because: (a) management decisions (antibiotic choice, fluid balance, renal dosing) are explicitly documented in discharge notes; (b) answering requires synthesizing across multiple patients: some with pneumonia alone, some with pneumonia + CKD, providing both the baseline and the modifier; (c) ground truth is derivable: chunks from admissions carrying both pneumonia and CKD ICD codes, plus chunks from pneumonia-only admissions for contrast. Each chunk contributes one patient’s management; the gold set is the collection of chunks whose union covers all required facets.

This query style aligns with the coverage hypothesis. Top- k will over-retrieve the most common pneumonia presentation (central cluster). **Coverage** ensures representation from CKD-modified management, elderly-specific complications (delirium, frailty), and atypical presentations.

3 MIMIC-IV Structured Data

3.1 Hosp Module

The **hosp** module captures hospital-wide data from admission to discharge. Key table groups (tables [highlighted in blue](#) are useful for us):

Patient tracking. Three tables describe demographics and movement:

- [patients](#): gender, [anchor_age](#) (deidentified age), date of death. Ages are shifted but intervals are preserved, so age *at admission* is recoverable.
- [admissions](#): admission/discharge times, admission type (emergency, elective, urgent), insurance, language, marital status, race, discharge location, in-hospital mortality flag.
- [transfers](#): ward-level ADT (admission–discharge–transfer) records tracking physical location throughout the stay.

Billing and diagnoses.

- [diagnoses_icd](#): up to 39 ICD-9-CM / ICD-10-CM codes per admission with approximate importance ordering ([seq_num](#)). Definitions in [d_icd_diagnoses](#). This is our primary source for condition-level labels and query scaffolding.
- [procedures_icd](#): billed procedures (ICD-9-PCS / ICD-10-PCS). Definitions in [d_icd_procedures](#).
- [drgcodes](#): Diagnosis Related Group codes (AP-DRG and MS-DRG) assigning an overall cost category to each hospitalization.
- [hcupsevents](#): HCPCS billing codes for provided services (e.g. mechanical ventilation, ICU care). Definitions in [d_hcpcs](#).

Measurements.

- **labevents**: laboratory results (chemistry, hematology, blood gases, etc.) with reference ranges and flags. Concept definitions in **d_labitems**. Multiple measurements may share a **specimen_id** (same blood draw).
- **microbiologyevents**: microbiology cultures with specimen type, organism, antibiotic sensitivities, and dilution results. Stored as a directed hierarchy (specimen → organism → antibiotic).
- **omr**: Online Medical Record outpatient measurements - blood pressure, height, weight, BMI, eGFR - available from both inpatient and outpatient visits.

Medications. Four tables track prescription and administration:

- **prescriptions + pharmacy**: provider orders with drug name, dose, route, and formulary details.
- **emar + emar_detail**: electronic Medicine Administration Record (barcode-scanned actual administrations). One row per administration event, linked to **pharmacy** and **poe** tables.

Administration. **services** (admitting service), **poe** and **poe_detail** (provider order entry - orders for imaging, consults, nutrition, medications).

3.2 ICU Module

The **icu** module (94,458 ICU stays, 65,366 unique patients) contains bedside data from the MetaVision clinical information system (iMDsoft), limited to patients admitted to an ICU. All events reference a **stay_id** (unique per ICU stay), defined in **icustays** with admission/discharge times and care unit names (e.g. MICU, SICU, CCU).

Event tables. Each documented item is an “event” grouped by data type:

Table	Content	Data Type
chartevents	Vital signs, GCS, vent settings, nursing assessments	Numeric time-series
inputevents	IV infusion rates/amounts (vasopressors, sedation)	Numeric
outputevents	Urine output, drain volumes	Numeric
procedureevents	Ventilation, CRRT, catheters (start/end times)	Categorical + timestamps
datetimeevents	Date-valued events (e.g., last dialysis)	Dates
ingredientevents	Ingredient-level breakdown of infusions	Numeric
d_items	Dimension table for itemid	Reference

Suitability for embeddings. ICU data is predominantly **numerical time-series** (vital signs, infusion rates, urine output) rather than natural language. These signals are not directly embeddable with sentence transformers, and converting them to text (e.g., “heart rate 92 bpm at 14:30”) would produce low-information chunks unlikely to match condition-level queries semantically. The ICU module is not a primary source for our corpus.

Metadata enrichment. ICU data could strengthen the **contextual embedding prefix** (Section 4.1) for patients who had ICU stays. The **mimic-code** repository already provides SQL concepts for severity scores (SOFA, SAPS-II, OASIS), ventilation status, and vasopressor use. A richer prefix:

Discharge note for elderly male, admitted for sepsis
(SOFA 14, mechanically ventilated, on vasopressors).
Comorbidities: CHF, CKD.
Section: Brief Hospital Course.

is more informative than one without ICU severity signals. However, this enrichment only matters for **ICU-relevant conditions** (sepsis, respiratory failure, shock, post-cardiac arrest) where severity

scores carry discriminative information. For routine admissions (e.g. elective surgery, uncomplicated pneumonia) it maybe doesn't make sense.

Part II

QA Dataset Generation Pipeline

4 Pipeline draft

4.1 Phases

Phase 1: Corpus Construction

1. **Select target conditions.** Query `diagnoses_icd` for ICD code groups with ≥ 200 (TBD based on dataset statistics. TODO: gather analytics) admissions.
Candidate conditions: stroke, myocardial infarction, pneumonia, sepsis, heart failure, COPD exacerbation, acute kidney injury, pulmonary embolism, diabetic ketoacidosis, cirrhosis, etc. Target: 50–100 conditions.
Prefer conditions with high **ICD co-occurrence richness**: conditions whose admissions carry diverse comorbidity profiles (e.g., pneumonia + CKD, pneumonia + CHF, pneumonia + diabetes) are more likely to produce multi-aspect candidate pools than conditions with uniform ICD profiles. This can be measured as the number of distinct ICD code sets across admissions for a given condition.
2. **Build per-condition candidate pools.** For each condition, gather all discharge note chunks from admissions with that ICD code. This is the full candidate set D for facility-location.
3. **Chunk by section + sub-problem.**
 - Split on section headers
 - Within `Brief Hospital Course`, split on # problem markers
4. **Attach metadata to each chunk:**
 - `subject_id`, `hadm_id`, section name
 - All ICD codes for that admission
 - Demographics: age group, gender, race, insurance
 - Admission type, discharge disposition
5. **Deduplicate same-patient chunks.** The same patient can have multiple hospitalizations with nearly identical boilerplate sections, e.g., `Past Medical History`, `Allergies`, and `Family History` are often copied verbatim across admissions.

Strategy: compute a content hash (e.g., SHA-256 of normalized text) per chunk and deduplicate by (`subject_id`, section, hash). Check also if it's worth considering *only* the most up to date note for a certain patient for a subset of repetitive sections (`Past Medical History`, etc.)

Phase 2: Embeddings Generation

Raw chunks are too patient-specific to match condition-level queries. Following the **Contextual Retrieval** approach [3], prepend a metadata prefix to each chunk before embedding:

```
[age_group] [gender] [race] patient admitted for [primary ICD description].  
Chief complaint: [chief_complaint].
```

Comorbidities: [top ICD descriptions].
Section: [section_name].

Check out Section 6 for more details.

The specific contents to include are to be decided, based on the queries we want to build.

Phase 3: Query Generation

1. Pick a primary condition (e.g., stroke, ICD-10: I63.*)
2. Pick 1–2 **modifier axes** from metadata. Each axis resolves to a filter on `hadm_id`, which is the join key to the discharge notes already chunked in Phase 1. Since chunk metadata was pre-attached in Phase 1 step 4, query generation simply selects which modifier values to combine into the prompt template.
 - **Comorbidity** (prior MI, diabetes, renal failure): co-occurring ICD codes on the same admission. A single `hadm_id` carries 5–20 ICD rows in `diagnoses_icd`; “stroke + diabetes” is the set of `hadm_ids` that have both I63.* and E11.*. The human-readable description for the prompt comes from `d_icd_diagnoses`.
 - **Demographic** (elderly, women, specific race): requires joining two tables. `patients` provides `gender` and `anchor_age`; `admissions` provides `race`, `insurance`, `marital_status`, and `admission_type`. Both join to `diagnoses_icd` via `subject_id` and `hadm_id` respectively.
 - **Complication** (post-surgical, recurrent, hemorrhagic transformation): draws from multiple sources. Co-occurring ICD codes capture complications (e.g., stroke I63.* + hemorrhage I61.* on the same `hadm_id`). `procedures_icd` identifies post-surgical modifiers. Recurrence is detected by finding the same `subject_id` with the target condition across multiple `hadm_ids`.

3. Prompt an LLM:

“Generate a clinical management question about [condition] that requires knowing how treatment or workup changes when [modifier] is present. The answer should require synthesizing management decisions from multiple patient cases, not textbook symptom lists.”

4. **Divergence pre-filter.** Before investing in expensive gold annotation, run facility-location and top- k on the candidate pool for the generated query and compute two cheap signals:
 - **Jaccard divergence:** if $J(S_{\text{cov}}, S_{\text{topk}}) \approx 1$, coverage selects nearly the same chunks as top- k $\Rightarrow$ skip this query, coverage adds nothing here. If $J \ll 1$, coverage diverges $\Rightarrow$ proceed.
 - **Facility-location gap:** $g(S_{\text{cov}}) - g(S_{\text{topk}})$ quantifies how much pool structure top- k ignores. Both S_{cov} and S_{topk} are already computed, so this is free.

Only queries where coverage meaningfully diverges from top- k proceed to gold annotation.

5. **Gold answer generation**

The idea: give an LLM access to the candidate pool, have it answer the query comprehensively, and for each piece of information in the answer **cite which chunk(s)** it drew from. The cited chunks become the **gold set**, each tagged with a facet label.

Facet annotation. The gold answer naturally covers several distinct information aspects — each one is a **facet**. Facet labels are not pre-defined; the LLM generates them during annotation. For each piece of information in its answer, the LLM cites which chunk(s) it drew from and assigns a short descriptive label for the aspect of care it addresses.

Example. For the query “How does management of acute stroke differ when the patient has concurrent CKD?”, the annotator might produce:

Facet label	What it captures	Example chunk content
anticoag_adjustment	CKD affects blood thinner dosing	Heparin dose reduced to 10u/kg/hr due to CrCl < 30
contrast_avoidance	CKD contraindicates contrast imaging	CT angiography deferred, nephropathy risk
bp_target	Different BP targets in stroke + CKD	Permissive HTN limit adjusted for renal function
dialysis_timing	When to initiate dialysis during acute stroke	Nephrology consulted for HD timing peri-stroke

Multiple chunks can map to the same facet (e.g., two different patients’ notes both discuss anticoagulation adjustment). Aspect recall does not require retrieving *all* chunks for a facet — just *at least one* per facet (the $\exists$ in the AR formula, Section 4.1).

This is where the coverage hypothesis applies: top- k might retrieve several chunks from the dominant facet (**anticoag_adjustment**, which appears in many notes) while missing niche facets (**dialysis_timing**, which appears in fewer notes and gets buried by similarity ranking). Coverage should recover more facets by pulling representatives from different regions of the pool.

Exhaustive tagging and strict evaluation. The map pass reads *every* chunk in the pool, so the LLM is instructed to tag **all** chunks that support each facet, not just the single best one. This makes the gold set comprehensive. At evaluation time, aspect recall uses **strict chunk-id matching**: a retrieved chunk contributes to a facet only if it appears in the gold set for that facet. No LLM judge is needed at evaluation time — the check is pure set intersection, making evaluation deterministic and reproducible.

Gold annotation via map-reduce. Since the corpus is heavily pruned before this stage (only conditions with the relevant ICD codes, only queries that pass the divergence filter), each per-condition candidate pool should not be too large. However, to face the long-context that the candidates may give rise to, we can:

- (a) **Map**: split the pool into batches of ~ 50 chunks ($\sim 10K$ tokens). For each batch, the LLM extracts facts relevant to the query, cites which chunk(s) each fact came from, and assigns a facet label to each citation.
- (b) **Reduce**: a merge pass synthesizes across batches, deduplicates facets, and produces the final gold answer with the complete (facet $\rightarrow$ chunk-ids) mapping.

6. (Optional) Query quality verifier

- A separate reader-LLM attempts to answer from baseline top- k chunks
- If it succeeds $\Rightarrow$ query does *not* require diversity $\Rightarrow$ discard
- If it fails but succeeds with gold chunks $\Rightarrow$ query requires diversity $\Rightarrow$ keep

Phase 4: Evaluation

- **Aspect recall**: fraction of annotated facets covered by selected k chunks (strict chunk-id matching).

$$AR(S) = |\{f \in F : S \cap G_f \neq \emptyset\}| / |F|$$
, where G_f is the set of gold chunk-ids tagged with facet f .
- **Fac score**: facility-location objective value $g(S) = \sum_{i \in D} \max_{j \in S} w_{ij}$
- **Jaccard**: overlap between method’s selection and top- k baseline
- **DocRec / FactRec**: gold document/fact fraction recovered

Methods to compare: top- k (baseline), MMR, gMMR, facility-location

4.2 Does Coverage Pull in Noise?

Facility-location maximizes $g(S) = \sum_{i \in D} \max_{j \in S} w_{ij}$ over the entire candidate pool D . It tries to *cover all of D* , including data points that are close to the query in embedding space but irrelevant to the query’s specific aspects. For example, a query about “pneumonia management in CKD (Complex Kidney Disease) patients” retrieves a candidate pool where most chunks are about pneumonia, some about CKD, and some are noise (superficially similar embeddings about unrelated respiratory complaints). Pure coverage would “waste” a selection on representing the noise region.

Mitigations:

1. **Candidate pool quality.** Coverage is applied as a *re-ranker* over N candidates retrieved by ANN (Section 4.3), not over the full corpus. If ANN retrieval is reasonable, most of D is query-relevant and noise is a small fraction with weak pull on the objective.
2. **λ tuning.** At $\lambda = 1$, coverage degrades to top- k ; at $\lambda = 0$, it ignores relevance entirely. The optimal λ is a thesis parameter. If noise is a problem, it manifests as coverage underperforming top- k at low λ , which we can detect in the experiments.

4.3 Computational Feasibility

Greedy facility-location is $O(|D| \cdot k)$ per selection step, so $O(|D| \cdot k^2)$ total. With $|D| \approx 3\text{M}$ chunks (331K notes $\times$ ~ 10 chunks each) and $k = 10$, this is $\sim 3 \times 10^8$ similarity computations per query, impractical.

Proposal: **two-stage pipeline**:

1. **Retrieval:** HNSW index retrieves $N \gg k$ candidates by approximate nearest neighbor search. Typical N : 100–1000. This is sub-linear in $|D|$.
2. **Re-ranking:** run facility-location (or MMR, gMMR, FPS) on the N candidates to select k .

N as a parameter. Too small: coverage has no room to diversify (all N candidates are near-duplicates from the dominant cluster). Too large: distant, irrelevant chunks enter the pool. N should be swept in experiments alongside λ and k .

5 LLM Usage and Data Use Agreement

Phases 3–4 require feeding MIMIC discharge note chunks into an LLM (query generation, gold answer annotation, optional quality verification). The PhysioNet Credentialed Health Data License 1.5.0 and DUA 1.5.0 impose two constraints:

1. **No sharing.** “*I will not share access to PhysioNet restricted data with anyone else.*” Sending clinical text to a commercial API arguably violates this clause.
2. **Security.** “*I will exercise all reasonable and prudent care to maintain the physical and electronic security of PhysioNet restricted data.*” Even with zero-retention API tiers, the data still transits and is processed on third-party servers outside our control.

Local inference. All LLM steps that touch patient data should run on a self-hosted model. Available hardware: NVIDIA L4 (23 GB VRAM). Speed? Performance of small LLMs?

Exception. Steps that use only *public* information (ICD code descriptions, condition names, modifier labels) can use commercial APIs, since no MIMIC text is transmitted. This applies to query template generation (Phase 3, step 3) but not to gold annotation.

6 Embeddings: considerations

6.1 Contextual Embeddings

A raw chunk from Brief Hospital Course might read:

“Pt had L-sided weakness, slurred speech, presented within 3h window. tPA administered. CT head negative for hemorrhage.”

Without context, embedding this chunk places it in a generic “neurological emergency” region. A query about “stroke symptoms” may not rank it highly, because the semantics is different.

Contextual Retrieval [3] Prepend contextual information to each chunk before embedding. Reported: 49% reduction in retrieval failure rate (combined with BM25 hybrid).

The original method uses an LLM to generate the prefix from the surrounding document, since document structure is unknown in the general case. MIMIC discharge notes, however, come with structured metadata (ICD codes, demographics, admission details) already parsed in relational tables, so we replace the LLM call with a deterministic template.

6.2 Chunking vs. Whole-Document Embedding

Long-context embedding models (e.g., BioClinical ModernBERT with 8K tokens, or Jina v2 with 8K tokens) can embed an entire discharge note as a single vector. Should we consider?

A whole-document embedding of a 2,000-token discharge note compresses heterogeneous content (chief complaint, clinical reasoning,...) into a single 768-dim vector. Mean-pooling over these diverse sections produces a vector that sits in a “generic admission” region of the embedding space, rather than near any specific clinical concept. This flattens the multi-cluster structure we need, so coverage brings no advantage.

When whole-document embedding? If queries target *patient-level* distinctions rather than *information-level* distinctions (e.g., “find patients similar to this case” rather than “what are the symptoms of X and how do they differ with Y”), whole-document embedding is more natural. This is based upon the queries we want to build, so maybe it’s good

6.3 Embedding Models for medical domain

- **BioLORD-2023 [10]**: sentence-transformer fine-tuned from `all-mpnet-base-v2`. 768-dim embeddings, 512-token context.
- **MedEmbed [4]**: family of sentence-transformers (small/base/large) fine-tuned from the BAAI/BGE series. The large variant produces 1024-dim embeddings, 512-token context. Apache-2.0 licence.
- **BioClinical ModernBERT Embeddings [11]**: sentence-transformer fine-tuned from BioClinical ModernBERT, a ModernBERT variant further pre-trained on PubMed abstracts and clinical text in the MIMIC lineage. 768-dim embeddings, up to 8,192-token context. Apache-2.0 licence.

7 Cluster Validation

Do candidate pools have the desired multi-cluster structure of Section 1?

7.1 Pipeline draft

1. **Embed** all chunks with `multi-qa-mpnet-base-cos-v1` (768 dimensions).

2. **UMAP** reduction to $\sim 15\text{--}30$ dimensions. UMAP [8] preserves local *and* global structure better than t-SNE.
3. **HDBSCAN** [5] with `min_cluster_size` tuned per condition.
4. **Validate** structure with:
 - **DBC**V [9]: density-based cluster validation, designed for HDBSCAN output
 - **Cluster size ratio**: $|C_{\max}|/|C_{\text{total}}|$ measures dominance of central cluster
 - **Density ratio**: $\rho_{\text{central}}/\rho_{\text{niche}}$
 - **Number of clusters + noise fraction**
5. **Cross-reference** HDBSCAN clusters with ICD sub-codes. If embedding clusters align with clinical sub-types, the structure is semantically meaningful.

7.2 Is Full Clustering Necessary?

The Jaccard divergence and facility-location gap signals used in Phase 3 step 4 already tell us whether coverage selects differently from top- k , without requiring any clustering. Full clustering (UMAP + HDBSCAN) is not needed as a pipeline gate, but remains useful as a post-hoc *explanatory* analysis: correlating cluster statistics (number of clusters, dominance ratio) with ΔAR can help explain *why* coverage helped on certain conditions. Ultimately, **aspect recall** is the real validation: if coverage achieves higher aspect recall than top- k on a condition, the pool structure was there.

7.3 High-Dimensional Clustering: Caveats

If we do use clustering for analysis, the following caveats apply.

Curse of dimensionality. HDBSCAN relies on point density. In high dimensions, pairwise distances concentrate [1] and density-based clustering becomes unreliable. This motivates the UMAP + HDBSCAN pipeline (used by BERTopic [7], Top2Vec [2]).

Is UMAP necessary? The concentration-of-distances argument applies to *uniformly distributed* high-dimensional data. Sentence embeddings are not uniform, they occupy a narrow anisotropic cone [6], and their intrinsic dimensionality is much lower than the ambient 768. HDBSCAN with cosine distance may find meaningful clusters directly in the original space *if* the effective dimensionality is low enough. This is an empirical question: run HDBSCAN both with and without UMAP and compare. If results agree, UMAP adds stochastic noise (sensitive to `n_neighbors`, `min_dist`) without benefit.

Anisotropy cuts both ways. Anisotropy means clusters may reflect frequency or style artifacts (all discharge notes share boilerplate language) rather than the semantic distinctions we care about. Cross-referencing HDBSCAN clusters with ICD sub-codes (Step 5 above) is essential: if clusters don't align with clinical sub-types, they are artifacts.

References

- [1] Charu C Aggarwal, Alexander Hinneburg, and Daniel A Keim. On the surprising behavior of distance metrics in high dimensional space. In *International conference on database theory*, pages 420–434. Springer, 2001.
- [2] Dimitar Angelov. Top2vec: Distributed representations of topics. *ArXiv*, abs/2008.09470, 2020.
- [3] Anthropic. Contextual retrieval, 2024. Blog post.
- [4] Abhinand Balachandran. Medembed: Medical-focused embedding models, 2024.

- [5] Ricardo J. G. B. Campello, Davoud Moulavi, and Joerg Sander. Density-based clustering based on hierarchical density estimates. In Jian Pei, Vincent S. Tseng, Longbing Cao, Hiroshi Motoda, and Guandong Xu, editors, *Advances in Knowledge Discovery and Data Mining*, pages 160–172. Springer Berlin Heidelberg, 2013.
- [6] Kawin Ethayarajh. How contextual are contextualized word representations? Comparing the geometry of BERT, ELMo, and GPT-2 embeddings. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- [7] Maarten Grootendorst. BERTopic: Neural topic modeling with a class-based TF-IDF procedure. *arXiv preprint arXiv:2203.05794*, 2022.
- [8] Leland McInnes, John Healy, Nathaniel Saul, and Lukas Großberger. Umap: Uniform manifold approximation and projection. *Journal of Open Source Software*, 3(29):861, 2018.
- [9] Davoud Moulavi, Pablo A Jaskowiak, Ricardo JGB Campello, Arthur Zimek, and Jörg Sander. Density-based clustering validation. In *Proceedings of the 2014 SIAM international conference on data mining*, pages 839–847. SIAM, 2014.
- [10] François Remy, Kris Demuynck, and Thomas Demeester. BioLORD-2023: Semantic textual representations fusing large language models and clinical knowledge graph insights. *Journal of the American Medical Informatics Association*, 31(9):1844–1855, February 2024.
- [11] Thomas Sounack, Joshua Davis, Brigitte Durieux, Antoine Chaffin, Tom J Pollard, Eric Lehman, Alistair EW Johnson, Matthew McDermott, Tristan Naumann, and Charlotta Lindvall. Bioclinical modernbert: A state-of-the-art long-context encoder for biomedical and clinical nlp. *arXiv preprint arXiv:2506.10896*, 2025.

Appendix

A Available Analytics from mimic-code

The official `mimic-code` repository provides SQL concepts, cohort definitions, and mapping tables that we can use.

Resource	Location	Role in Our Pipeline
Charlson Comorbidity Index	<code>mimic-iv/concepts/\protect\penalty\z@comorbidity/charlson.sql</code>	Maps ICD-9/10 codes to 16 clinical categories (MI, CHF, stroke, COPD, diabetes, renal, cancer...). Primary tool for condition selection and comorbidity modifier axes .
CCS hierarchical grouping	<code>mimic-iii/concepts_postgres/\protect\penalty\z@diagnosis/ccs_dx.sql</code> + <code>ccs_multi_dx.csv.gz</code>	4-level hierarchical ICD-9 → clinical category mapping. Useful for building the macrocategory index in hierarchical retrieval (Phase 2).
Concept maps	<code>mimic-iv/concepts/\protect\penalty\z@concept_map/</code>	ICD → LOINC/OMOP/SNOMED CSVs. Cross-referencing structured data when building contextual embedding prefixes.
Age / demographics	<code>mimic-iv/concepts/\protect\penalty\z@demographics/age.sql</code>	Age cohort extraction. Needed for demographic modifier axes.
Severity scores	<code>mimic-iv/concepts/\protect\penalty\z@score/</code> (SOFA, SAPS-II, LODS, APS-III, OASIS)	Patient acuity signals. Could serve as additional metadata for contextual embeddings or as a proxy for clinical complexity.